

@holokai/neuron-sdk

SDK for building neurons — processes that connect to the BigBrain gateway and execute tasks on behalf of workflows.

A neuron advertises one or more **capabilities** (typed task handlers), opens a persistent connection to the gateway, and receives work via a pluggable `Transport`. Two transports ship in the SDK:

- **HttpTransport** — SSE for inbound, HTTPS POST for outbound. The default for any neuron talking to a remote BigBrain gateway. Re-exported from the package root.
- **IpcTransport** — Electron renderer ↔ main IPC. For neurons running inside an app that embeds BigBrain in main; no network hop, no JWT in the renderer. Imported from the `/ipc` subpath.

The gateway routes tasks to the right neuron based on capability type and scope. The built-in `apps/worker` uses this SDK over HTTP; the in-process Holokai Desktop renderer uses it over IPC; both share the same `Neuron` class and handler API.

Installation

```
pnpm add @holokai/neuron-sdk
# or: npm install @holokai/neuron-sdk
# or: yarn add @holokai/neuron-sdk
```

Runs in **Node 20+** (uses the global `fetch` and `ReadableStream`) and **modern browsers** (Chromium, Firefox, Safari latest). The package has no `node:*` imports — Holokai Desktop / Electron and the in-browser neuron-tester both consume the same SDK build.

Gateway compatibility (protocol 1.1.0, HOL-2926). Since protocol 1.1.0 the SDK advertises each capability's contract (`description` + input/output JSON Schemas) on the register frame. The gateway validates register frames **strictly**, so a gateway older than HOL-2926 rejects them with `INVALID_FRAME` — **deploy the gateway before publishing or rolling out any 1.1.0-protocol SDK (including neuron-tester OTA)**. A rejected registration leaves the SSE stream open but the neuron lease-less; the SDK surfaces it via the `error` event (`registration rejected by gateway: ...`) so apps can react.

Quick start

```
import { Neuron } from '@holokai/neuron-sdk';

const neuron = new Neuron({
  gatewayUrl: 'https://bigbrain.holokai.dev',
  neuronId: 'my-worker-001', // stable per process/device
  auth: () => process.env.GATEWAY_TOKEN!, // or an async refresh function
});

neuron.handle(
  { type: 'core/script.python', scope: 'any', concurrency: 4 },
  async (input, ctx) => {
    ctx.log.info({}, 'running python script');
    const result = await runPython(input as { code: string }, ctx.signal);
    return { output: result };
  },
);

await neuron.start();

process.on('SIGTERM', () => neuron.stop({ drain: true, timeoutMs: 30_000 }));
```

Authentication

The `auth` callback supplies the bearer token for every POST and SSE open. How you obtain tokens depends on where the neuron runs (full model: `docs/NEURON_AUTH_SPEC.md` in the bigbrain repo):

- **Desktop-embedded neurons** run under the signed-in user: exchange the user's Moku JWT for an `aud=bigbrain` token (RFC 8693) and return it from `auth`.
- **Headless / enterprise neurons** enroll once via Moku's device flow, then mint short-lived tokens from the stored client credential — no static long-lived JWT anywhere.

Enroll (one-time, headless neurons)

```
npx @holokai/neuron-sdk enroll --moku-url https://moku.example \
  --name warehouse-connector --capabilities org/acme/erp.read,org/acme/erp.write
```

Prints a verification URL + user code; an org admin approves in Moku. The client credential is written to `~/.holokai/neuron-credentials.json` (0600; override with `--credentials-file`). The requested capabilities are display hints for the approval screen — what the neuron may actually advertise is governed by the org's capability catalog. Polling is resilient: transient server errors (409/5xx) and network blips are retried until the device code expires; only a denial, expiry, or protocol error is terminal.

Also available programmatically: `enroll(opts)` from `@holokai/neuron-sdk/auth`.

Steady state

```
import { Neuron } from '@holokai/neuron-sdk';
import { createClientCredentialsAuth, defaultCredentialPath } from '@holokai/neuron-sdk/auth'

const neuron = new Neuron({
  gatewayUrl: 'https://bigbrain.holokai.dev',
  neuronId: 'warehouse-connector-001',
  auth: createClientCredentialsAuth(defaultCredentialPath()),
});
```

`createClientCredentialsAuth` mints `aud=bigbrain` tokens via the OAuth2 `client_credentials` grant, caches them against the response's `expires_in` (proactive refresh, single-flight), and re-reads the credential file on `invalid_client` so an operator-rotated secret is picked up without a restart. If the credential was revoked in Moku it throws `CredentialRevokedError` — re-enroll to recover.

`@holokai/neuron-sdk/auth` is **Node-only** (it reads the credential file); the root and `./ipc` entries remain browser-clean. Python SDK parity ships in `holokai-neuron-sdk` `≥ 0.2.0` (`holokai_neuron_sdk.auth: holokai-neuron-sdk enroll + create_client_credentials_auth`) — same credential file, interchangeable across both SDKs.

Core concepts

Capability

A capability is a named contract a neuron can fulfill. It has:

Field	Type	Description
<code>type</code>	<code>string</code>	Namespaced identifier: <code>{owner}/{package}[/{name}][.{variant}]</code>
<code>scope</code>	<code>'any' { onBehalfOf: userId }</code>	Who this handler serves
<code>concurrency</code>	<code>number</code>	Max parallel tasks for this capability

Examples:

- `core/script.python` — server-side Python execution, any user
- `desktop/mcp/drive.read` — user's Google Drive, scoped to that user
- `org/01HBQ8ZK/custom-tool` — org-specific capability

Namespace authorization

The gateway enforces which identities may advertise which namespaces:

Namespace prefix	Permitted advertisers	Required scope
core/*	Server service identities	any
desktop/*	Attested desktop binaries	{ onBehalfOf: userId }
org/{orgId}/*	Tenant-installed extensions	any within org
user/{userId}/*	Personal extensions	{ onBehalfOf: userId }

Registration is rejected (HTTP 403) if the JWT is not authorized for the namespace.

Scope

- **'any'** — the neuron handles tasks for any user. Shared queue; no per-user binding.
- **{ onBehalfOf: userId }** — the neuron handles tasks for that specific user only. The JWT's user binding must match; the gateway enforces this at registration time.

One neuron may advertise mixed-scope capabilities:

```
neuron.handle({ type: 'desktop/mcp/drive.read', scope: { onBehalfOf: 'user:abc' }, concurrency: 1 })
neuron.handle({ type: 'core/script.python', scope: 'any', concurrency: 1 })
```

Lease

A lease is how the gateway delivers a task. When the broker dispatches a task to a neuron, the gateway:

1. Writes a DB row with `leaseId` and expiry.
2. Sends a `lease` frame over SSE.

The SDK validates input, invokes your handler, and sends either `ack` (success) or `nack` (failure) over HTTPS POST. Until the ack arrives, the message remains unacknowledged in the broker — if the neuron disconnects, the broker redelivers automatically.

Handlers must be idempotent. A task that completed a side effect but failed to ack will be retried. Use capability-native idempotency keys (e.g., Google Drive `requestId`, HTTP `Idempotency-Key`).

Neuron ID

`neuronId` must be **stable across restarts** for the same process/device. The gateway uses it to correlate reconnects. Pick something meaningful and persistent:

```
// Server process
neuronId: `worker-${process.env.HOSTNAME}-${process.pid}`

// Desktop app – persist a UUID to disk
neuronId: await readOrCreateStoredId('/var/lib/myapp/neuron-id')
```

API reference

`new Neuron(opts: NeuronOptions)`

`NeuronOptions` is a discriminated union — pick the convenience shape (HTTP) or pass a transport explicitly.

```
type NeuronOptions = NeuronOptionsWithHttp | NeuronOptionsWithTransport;

interface NeuronCommonOptions {
  neuronId: string; // Stable per process/device
  logger?: NeuronLogger; // Plug in pino, console, etc. Default: silent
  heartbeatIntervalMs?: number; // Default 10_000 (10s)
}

// Convenience: SDK builds an HttpTransport for you. Most common path.
interface NeuronOptionsWithHttp extends NeuronCommonOptions {
  gatewayUrl: string; // Base URL, e.g. 'https://bigbrain.holokai.dev'
  auth: AuthCallback; // Token provider (see below)
  initialReconnectDelayMs?: number; // Default 500ms
  maxReconnectDelayMs?: number; // Default 30_000 (30s)
  fetch?: typeof fetch; // Override for tests
}

// Power-user: bring your own Transport. Used for IpcTransport, custom
// transports, or HttpTransport instances built explicitly.
interface NeuronOptionsWithTransport extends NeuronCommonOptions {
  transport: Transport;
}
```

See the [Transports](#) section for `HttpTransport`, `IpcTransport`, and the `Transport` interface.

`neuron.handle(reg, handler)`

Register a handler for a capability. Must be called **before** `neuron.start()`.

```
neuron.handle(reg: CapabilityRegistration, handler: Handler): void
```

```
interface CapabilityRegistration {
  type: string; // Namespaced capability type
  scope: 'any' | { onBehalfOf: string };
  concurrency: number; // Max parallel executions
  inputSchema?: JsonSchema; // Optional; validated before invoking handler
  outputSchema?: JsonSchema; // Optional; validated before acking
}

type Handler<TInput = unknown, TOutput = unknown> = (
  input: TInput,
  ctx: HandlerContext,
) => Promise<TOutput> | TOutput;
```

If `inputSchema` is provided and the lease input fails validation, the handler is never called — the SDK sends a `schema-mismatch` nack automatically. Same for `outputSchema` on the return value.

Schema precedence. For each of input and output the SDK picks exactly one validator:

1. The zod schema from `defineCapability`, when the capability was registered that way. JSON Schema can't express `.refine()`, `.transform()`, or `.preprocess()`, so the zod contract stays authoritative and its transforms are what the handler sees and what goes on the wire.
2. Otherwise the JSON Schema passed here on `CapabilityRegistration`.

A capability registered with neither is not validated client-side — its handler runs on whatever arrives.

The `inputSchema` / `outputSchema` that ride on the lease frame are **not** consulted. They can only ever be your own schemas handed back to you — the register frame carries no schemas, so the server's only source is the `POST /plan` handler catalog the client supplies. And the planner doesn't persist those onto task rows, so on real workflow tasks the lease frame's schemas are absent entirely; the only case where they arrive populated is the dev-only `/__dev__/inject-task` route used by `neuron-tester`, which echoes back what the capability registered.

This is the only validation that runs. The gateway's ack-time check validates against that same (empty) `lease.outputSchema`, so it does not currently catch anything. If a capability declares no schema here, its input and output are unchecked end to end.

Note: the Go and Python SDKs currently prefer the lease frame's schema and fall back to the registered one. TypeScript deliberately diverges here; the other two are expected to follow.

`neuron.start()`

Opens the SSE stream, posts `register` with all advertised capabilities, and starts the heartbeat loop.

```
await neuron.start(): Promise<void>
```

Throws if no handlers have been registered. Subsequent calls are no-ops.

`neuron.stop(opts?)`

Graceful shutdown: refuse new leases, drain in-flight handlers, disconnect.

```
await neuron.stop({ drain?: boolean, timeoutMs?: number }): Promise<void>
```

Option	Default	Description
<code>drain</code>	<code>true</code>	Wait for in-flight handlers to finish
<code>timeoutMs</code>	<code>30_000</code>	After this, abort remaining in-flight tasks

Queued leases (received but not yet started due to concurrency limits) are nacked `retryable` on shutdown so the broker redelivers immediately.

`neuron.updateCapability(reg)`

Change concurrency for an already-registered capability at runtime. Useful when available resources change (e.g., desktop plugged into power).

```
await neuron.updateCapability({ type: string, concurrency: number }): Promise<void>
```

Notifies the gateway via `update-capability` so it can tune AMQP prefetch. Safe to call while tasks are in flight — in-flight tasks finish at the old concurrency; new slots open at the new limit.

`HandlerContext`

The second argument to every handler:

```
interface HandlerContext {  
  task: TaskPayload; // Full task metadata  
  leaseId: string; // Gateway lease ID (for logging/debugging)  
  signal: AbortSignal; // Fires on gateway cancel or SDK shutdown  
  log: NeuronLogger; // Logger scoped to this lease  
  progress(update: { percent?: number; message?: string; data?: unknown }): void;  
  fail(reason: string, opts?: { code?: string }): never;  
}
```

`ctx.signal`

Pass this to any fetch/HTTP calls so they abort cleanly:

```
const resp = await fetch(url, { signal: ctx.signal });
```

When the gateway sends a `cancel` frame (workflow was cancelled by the user), `ctx.signal` fires and your handler should throw/return promptly. An `AbortError` from a cancelled signal is treated as `nack(kind='cancelled')` — benign, not counted as a failure.

`ctx.progress()`

Report progress for long-running tasks. Non-blocking; safe to call many times.

```
ctx.progress({ percent: 50, message: 'halfway done' });
ctx.progress({ data: { rowsProcessed: 1200 } });
```

`ctx.fail()`

Terminal failure — nacks the lease with `kind='terminal'` (no retry). Use when the task can never succeed regardless of retries.

```
if (!validApiKey) {
  ctx.fail('invalid API key', { code: 'INVALID_API_KEY' });
  // unreachable — fail() throws internally
}
```

AuthCallback

```
type AuthCallback = () => Promise<string> | string;
```

Called on every POST and SSE open. The SDK retries once on 401, so implement token refresh inside the callback:

```
let token = await fetchInitialToken();

const auth = async () => {
  if (isExpired(token)) {
    token = await refreshToken();
  }
  return token.accessToken;
};
```

NeuronLogger

```
interface NeuronLogger {
  debug(meta: Record<string, unknown>, msg?: string): void;
  info(meta: Record<string, unknown>, msg?: string): void;
  warn(meta: Record<string, unknown>, msg?: string): void;
  error(meta: Record<string, unknown>, msg?: string): void;
}
```

Matches pino's call signature directly. To use pino:

```
import pino from 'pino';
const logger = pino({ level: 'info' });
const neuron = new Neuron({ ..., logger });
```

defineCapability()

A typed helper that produces a capability definition from zod schemas, with handler inputs/outputs fully inferred:

```
import { defineCapability, Neuron } from '@holokai/neuron-sdk';
import { z } from 'zod';

const pythonCap = defineCapability({
  type: 'core/script.python',
  description: 'Run a Python script in a sandbox',
  input: z.object({ code: z.string(), args: z.record(z.unknown()).optional() }),
  output: z.object({ stdout: z.string(), exitCode: z.number() }),
  scope: 'any', // optional default; overridable at handle() time
  concurrency: 4, // optional default; overridable at handle() time
});

// Handler input/output inferred – no manual type annotations:
neuron.handle(pythonCap, async (input, ctx) => {
  // input: { code: string; args?: Record<string, unknown> }
  return { stdout: '...', exitCode: 0 };
});

// Or override defaults per deployment:
neuron.handle(pythonCap, { scope: { onBehalfOf: 'user-123' }, concurrency: 8 }, fn);
```

Validation uses the **original zod schemas** (so `.refine()`, `.transform()`, `.preprocess()`, `.brand()` all run as expected). The JSON Schema produced by `zod-to-json-schema` is shipped on the wire to the gateway and planner but is not used as the runtime validator — that avoids silently dropping zod features that don't translate.

The fall-back `neuron.handle(registration, fn)` shape (raw JSON Schema + capability fields) is still accepted for low-level callers that don't have zod schemas on hand.

Events

Neuron extends `EventEmitter` from `eventemitter3` — a browser-clean, type-safe drop-in for Node's `EventEmitter`. Subscribe with `.on()`, `.once()`, or `.off()`:

```
const off = neuron.on('frame', (frame) => {
  log.append(frame); // see every parsed SSE frame
});
off(); // unsubscribe

neuron.once('connection:registered', ({ sessionId }) => {
  console.log(`registered as ${sessionId}`);
});

neuron.on('connection:reconnecting', ({ attempt, delayMs, reason }) => {
  statusBar.show(`reconnecting (attempt ${attempt}, ${delayMs}ms)`, reason);
});
```

The full event map:

Event	Payload	When
<code>connection:open</code>	<i>none</i>	SSE handshake succeeded; before <code>register</code> is posted
<code>connection:registered</code>	<code>{ sessionId: string; capabilities: number }</code>	Gateway accepted the register frame
<code>connection:close</code>	<code>reason: string</code>	SSE stream ended (clean shutdown or transient error)
<code>connection:reconnecting</code>	<code>{ attempt, delayMs, reason }</code>	Reconnect scheduled, before backoff delay
<code>error</code>	<code>err: Error</code>	Transport / parse error (reconnect handling is automatic)
<code>frame</code>	<code>frame: ServerToNeuronFrame</code>	Pass-through of every parsed inbound SSE frame

The `frame` event is a firehose mirror of the gateway's `/neuron/events` SSE stream. It fires for **every** frame including ones the SDK handles internally (`lease`, `cancel`) — useful for debug UIs, observability, and audit logs. The high-level `neuron.handle()` API is implemented as a `frame` listener filtered to `lease` frames, so `handle()` and direct `frame` subscribers compose without conflict.

Subscribe **before** calling `neuron.start()` if you care about events that fire as part of the connection sequence (`connection:open`, the initial `registered` frame). Late subscribers will still see subsequent events.

The event surface is designed to grow — future entries (e.g. `'workflow:task-completed'` for cross-neuron notifications) slot in without breaking subscribers.

Transports

A `Transport` is the wire-level abstraction the rest of the SDK builds on. It carries typed frames in both directions: inbound `ServerToNeuronFrame`s via callbacks, outbound `ClientFrame`s via `post(...)`. The `Neuron` class only ever talks to a `Transport` — adding a new transport (e.g. WebSocket, native messaging) means implementing one interface.

```
interface Transport {
  subscribe(callbacks: TransportCallbacks): void;
  start(): Promise<void>;
  stop(): Promise<void>;
  post<T extends ClientFrame>(path: PostPath, frame: T): Promise<unknown>;
  isConnected(): boolean;
}

interface TransportCallbacks {
  onOpen?: () => void;
  onClose?: (reason: string) => void;
  onFrame?: (frame: ServerToNeuronFrame) => void;
  onError?: (err: Error) => void;
  onReconnecting?: (info: { attempt: number; delayMs: number; reason: string }) => void;
}
```

Both `Transport` and `TransportCallbacks` are exported from the package root.

HttpTransport

Default transport. SSE for inbound, HTTPS POST for outbound. Reconnects with exponential backoff (default: 500ms → 30s ceiling). Refreshes the auth token via your `AuthCallback` and retries once on 401.

```
import { Neuron, HttpTransport } from '@holokai/neuron-sdk';

const transport = new HttpTransport({
  gatewayUrl: 'https://bigbrain.holokai.dev',
  neuronId: 'my-worker-001',
  auth: () => process.env.GATEWAY_TOKEN!,
  // optional:
  initialReconnectDelayMs: 500,
  maxReconnectDelayMs: 30_000,
});

const neuron = new Neuron({ neuronId: 'my-worker-001', transport });
```

In practice you rarely construct `HttpTransport` directly — passing `gatewayUrl` + `auth` to `Neuron` builds it for you. Construct explicitly when you want to share a fetch implementation, swap the auth strategy, or instrument the transport in tests.

```
interface HttpTransportOptions {
  gatewayUrl: string;
  neuronId: string;
  auth: AuthCallback;
  logger?: NeuronLogger;
  initialReconnectDelayMs?: number; // default 500ms
  maxReconnectDelayMs?: number;    // default 30_000ms
  fetch?: typeof fetch;
}
```

Ipctransport

For renderer-process neurons running inside an Electron app that hosts BigBrain in the main process. Leases arrive over Electron IPC, not over the network. Imported from the `/ipc` subpath so non-Electron consumers don't pay for it:

```
import { Neuron } from '@holokai/neuron-sdk';
import { Ipctransport } from '@holokai/neuron-sdk/ipc';

// `bridge` comes from the renderer's preload — see below.
const transport = new Ipctransport({ bridge: window.bigbrainNeuron, neuronId });
const neuron = new Neuron({ neuronId, transport });

neuron.handle({ type: 'desktop/local-tool', scope: { onBehalfOf: userId }, concurrency: 1 },
await neuron.start());
```

Differences from `HttpTransport` :

- **No reconnect logic** — IPC has no transient socket. If the bridge throws, the failure surfaces via `onError` and the consumer decides what to do.
- **No auth callback** — the renderer never sees the JWT. The Electron main process attaches `Authorization` headers when it forwards POSTs to the embedded Fastify gateway.
- **No keep-alive** — IPC channels are process-local; nothing to time out.

The SDK does **not** depend on Electron. The renderer's preload script is responsible for constructing a concrete bridge that satisfies the `NeuronIpcBridge` shape and exposing it via `contextBridge` :

```
interface NeuronIpcBridge {
  attach(neuronId: string): Promise<{ channel: string }>;
  detach(neuronId: string): Promise<void>;
  subscribe(channel: string, cb: (frame: ServerToNeuronFrame) => void): () => void;
  invoke<T extends ClientFrame>(path: PostPath, frame: T): Promise<unknown>;
}
```

Typical preload wiring (Electron-side, app code — not part of the SDK):

```
// preload.ts
import { contextBridge, ipcRenderer } from 'electron';

contextBridge.exposeInMainWorld('bigbrainNeuron', {
  attach: (neuronId) => ipcRenderer.invoke('bigbrain:attach', neuronId),
  detach: (neuronId) => ipcRenderer.invoke('bigbrain:detach', neuronId),
  subscribe: (channel, cb) => {
    const listener = (_e, frame) => cb(frame);
    ipcRenderer.on(channel, listener);
    return () => ipcRenderer.off(channel, listener);
  },
  invoke: (path, frame) => ipcRenderer.invoke('bigbrain:post', path, frame),
});
```

```
interface IpcTransportOptions {
  bridge: NeuronIpcBridge;
  neuronId: string;
  logger?: NeuronLogger;
}
```

Examples

Simple any-scoped capability

```
import { Neuron } from '@holokai/neuron-sdk';

const neuron = new Neuron({
  gatewayUrl: process.env.GATEWAY_URL!,
  neuronId: `worker-${process.env.HOSTNAME}`,
  auth: () => process.env.GATEWAY_TOKEN!,
});

neuron.handle(
  { type: 'core/http', scope: 'any', concurrency: 8 },
  async (input, ctx) => {
    const { url, method = 'GET', headers = {}, body } = input as {
      url: string;
      method?: string;
      headers?: Record<string, string>;
      body?: string;
    };

    const resp = await fetch(url, { method, headers, body, signal: ctx.signal });
    const text = await resp.text();

    return { status: resp.status, body: text };
  },
);

await neuron.start();
```

User-scoped capability (desktop / MCP)

```
const userId = await getUserIdFromToken(authToken);

neuron.handle(
  { type: 'desktop/mcp/drive.read', scope: { onBehalfOf: userId }, concurrency: 1 },
  async (input, ctx) => {
    const { fileId } = input as { fileId: string };
    ctx.progress({ message: 'fetching file metadata' });

    const file = await driveClient.files.get({ fileId, signal: ctx.signal });
    ctx.progress({ percent: 50, message: 'downloading content' });

    const content = await driveClient.files.download({ fileId, signal: ctx.signal });
    return { name: file.name, mimeType: file.mimeType, content };
  },
);
```

Long-running task with progress

```
neuron.handle(  
  { type: 'core/data.transform', scope: 'any', concurrency: 2 },  
  async (input, ctx) => {  
    const { rows } = input as { rows: unknown[] };  
    const results: unknown[] = [];  
  
    for (let i = 0; i < rows.length; i++) {  
      if (ctx.signal.aborted) break; // respect cancellation  
  
      results.push(await processRow(rows[i]));  
      ctx.progress({ percent: Math.round(((i + 1) / rows.length) * 100) });  
    }  
  
    return { results };  
  },  
);
```

Error handling

```
neuron.handle(  
  { type: 'core/llm', scope: 'any', concurrency: 4 },  
  async (input, ctx) => {  
    const { prompt, model } = input as { prompt: string; model: string };  
  
    let response;  
    try {  
      response = await llmClient.complete({ prompt, model, signal: ctx.signal });  
    } catch (err) {  
      if ((err as Error).name === 'AbortError') {  
        throw err; // let the SDK handle cancellation  
      }  
      if ((err as any).status === 400) {  
        // Bad request – no point retrying  
        ctx.fail(`invalid LLM request: ${err as Error}.message`, { code: 'INVALID_REQUEST' })  
      }  
      // Anything else: throw to get a retryable nack  
      throw err;  
    }  
  
    return { text: response.text, tokensUsed: response.usage.total };  
  },  
);
```

Graceful shutdown

```
await neuron.start();

const shutdown = async (signal: string) => {
  console.log(`${signal} received - draining...`);
  await neuron.stop({ drain: true, timeoutMs: 30_000 });
  process.exit(0);
};

process.on('SIGTERM', () => shutdown('SIGTERM'));
process.on('SIGINT', () => shutdown('SIGINT'));
```

Runtime concurrency update

```
import { powerMonitor } from 'electron'; // or any power/resource API

powerMonitor.on('on-ac', async () => {
  await neuron.updateCapability({ type: 'core/script.python', concurrency: 8 });
});

powerMonitor.on('on-battery', async () => {
  await neuron.updateCapability({ type: 'core/script.python', concurrency: 1 });
});
```

Configuration reference

Option	Default	Applies to	Notes
<code>neuronId</code>	—	both	Stable per device/process; reused on reconnect
<code>logger</code>	silent no-op	both	Plug in pino or console
<code>heartbeatIntervalMs</code>	<code>10_000</code>	both	Sent every 10s; renews lease expiries
<code>transport</code>	—	with-transport shape	Pre-built <code>Transport</code> (e.g. <code>IpcTransport</code> or a custom <code>HttpTransport</code>)
<code>gatewayUrl</code>	—	with-http shape	Base URL, no trailing slash
<code>auth</code>	—	with-http shape	Called on every POST + SSE open
<code>initialReconnectDelayMs</code>	<code>500</code>	with-http shape	First reconnect delay (HTTP only — IPC doesn't reconnect)
<code>maxReconnectDelayMs</code>	<code>30_000</code>	with-http shape	Reconnect delay ceiling (exponential backoff)
<code>fetch</code>	<code>global fetch</code>	with-http shape	Override for tests

Protocol details

This section covers enough to debug; the wire-level frame schemas are bundled into the SDK and exported alongside the public types.

Transport

The SDK abstracts the wire as a `Transport` (see [Transports](#) above). The frame types below are the same regardless of transport — only the carrier differs.

`HttpTransport` (default):

- **Server** → **Neuron**: SSE stream at `GET /neuron/events?neuronId={id}`
- **Neuron** → **Server**: HTTPS POST endpoints (one per frame type)
- No WebSockets — SSE + HTTPS POST traverses enterprise proxies cleanly.

`IpcTransport` (Electron renderer):

- **Server** → **Neuron**: per-neuron Electron IPC channel obtained from `bridge.attach()`.
- **Neuron** → **Server**: `bridge.invoke(path, frame)`, which the main process forwards to the embedded Fastify gateway via `app.inject(...)` — same auth and validation chain as HTTP.
- No reconnect, no auth callback, no keep-alive.

Frame types

Neuron → Server (POST):

Frame	Path	When
register	/neuron/register	On SSE open
update-capability	/neuron/update-capability	On concurrency change
heartbeat	/neuron/heartbeat	Every 10s while connected
ack	/neuron/ack	Handler succeeded
nack	/neuron/nack	Handler failed
progress	/neuron/progress	During handler execution

Server → Neuron (SSE):

Frame	When
registered	Gateway accepted the register frame; effective capabilities returned
lease	Task to execute
cancel	Gateway cancels a specific lease (workflow cancelled)
capability-revoked	A capability was deauthorized mid-session
disconnect	Gateway is closing the connection cleanly


Heartbeat

Every ~10s the SDK posts:

```
{
  "type": "heartbeat",
  "neuronId": "worker-host-1234",
  "inFlight": [
    { "leaseId": "01HB...", "taskId": "task-abc", "startedAt": "2024-01-01T12:00:00Z" }
  ],
  "sentAt": "2024-01-01T12:00:10Z"
}
```

The gateway renews lease expiries for every `leaseId` in `inFlight` in a single DB round-trip. **Only running leases appear in `inFlight`** — queued leases (admitted but waiting for a concurrency slot) are excluded so the reaper doesn't mistakenly extend their TTL.

Lease lifecycle



Reconnection

The SDK uses exponential backoff (default: 500ms initial, doubling, 30s ceiling). On reconnect:

1. SSE stream re-opens.
2. `register` is posted with the full capability list.
3. Heartbeat loop resumes.

In-flight leases from before the disconnect were never acked — the broker redelivers them. The reaper flips stale `assigned` DB rows back to `pending`. Your handlers will receive the same task again; make them idempotent.

Troubleshooting

`Neuron.handle()` must be called before `start()`

Register all capabilities before calling `start()`. For runtime changes, use `updateCapability()`.

`Neuron.start()` requires at least one handler

You must call `neuron.handle(...)` at least once before `start()`.

Registration rejected with HTTP 403

Your JWT is not authorized for the namespace you're trying to advertise. Check:

- The capability type uses the correct namespace prefix (`core/*` requires a service identity, not a user JWT).
- For user-scoped capabilities, the `onBehalfOf` value matches the user in your JWT.

Tasks not being received

1. Check the `registered` log line — it shows the **effective** capabilities after the gateway's policy/auth filters. If a capability is missing, the gateway filtered it.
2. Check `neuronId` — if it changes across restarts, the gateway may have stale session state.
3. Check heartbeat logs — if heartbeats are failing, the lease reaper may be reclaiming tasks faster than they're running.

Handler keeps getting re-delivered

The handler completed but the `ack` POST failed (gateway returned 4xx or 5xx). Look for `ack POST rejected by gateway` in logs. Root causes:

- Output failed the `outputSchema` check — fix your handler's return shape.
- Gateway was restarting — the next delivery will succeed once it's back.

Tasks stuck in `assigned` state

Heartbeats aren't reaching the gateway, so the reaper hasn't reclaimed them. Check:

- Is the SSE stream still connected? (`transport reconnecting` log entries)
- Are heartbeat POSTs failing? (`heartbeat post failed` log entries)
- Is `heartbeatIntervalMs` misconfigured to a very large value?

Auth 401 on every request

The `auth` callback isn't returning a valid token. If using a refresh flow, ensure the callback fetches a fresh token on each call (the SDK caches nothing — it's your callback's job).

Releasing (maintainers)

`@holokai/neuron-sdk` is published to public npm by BigBrain's Release workflow (`.github/workflows/release.yml`). There is no npm token on anyone's laptop and none stored in the repo, and pushing a tag publishes nothing.

Pick the version

Follow semver:

- Pilot phase, we are in `0.x` . Because the protocol is still absorbing pilot-developer feedback, a **breaking** change bumps the **minor** (`0.2.0` to `0.3.0`); a backward-compatible change or fix bumps the **patch** (`0.2.0` to `0.2.1`). This is the standard `0.x` convention and is what buys us room to break things on feedback without a major bump.

- Hold `1.0.0` until the public API has survived pilot feedback and we're ready to commit to "breaking changes bump the major." Jumping to `1.0.0` is a stability promise, not a milestone. Don't spend it early.

Cut a release

1. Open **Actions, Release, Run workflow** on branch `main`.
2. Fill `neuron-sdk` with the new version (`X.Y.Z`). Leave the other lines (base, embedded, synapse) empty unless they are being released too; an empty line is not touched. Tick **plan only** to read the plan without cutting anything.
3. The plan job checks the version (shaped `X.Y.Z`, not older than `packages/neuron-sdk/package.json`, tag `neuron-sdk-vX.Y.Z` not taken) and writes the plan to the run summary.
4. A reviewer of the `release` environment approves the run. That is the only approval.
5. The cut writes the version into `packages/neuron-sdk/package.json`, commits it to `main`, tags `neuron-sdk-vX.Y.Z` and pushes both. If the file already carries the version (a bump merged by pull request), it only tags.
6. From the commit the cut pushed, the run builds and tests the SDK (the `prepublish turbo` task), packs it, and publishes the tarball to npm from a GitHub-hosted runner. On success the new version is live at [@holokai/neuron-sdk](https://www.npmjs.com/package/@holokai/neuron-sdk).

The tag is prefixed `neuron-sdk-v` (the Python SDK uses `neuron-sdk-python-v`) so the two never collide. The workflow file's header comments describe every job and how to recover a failed run.

No `--provenance`: bigbrain is a private repo, so npm's provenance attestation (a public transparency-log link back to the source) adds little and has historically errored from private repos. Trusted-publishing auth is unaffected. Add the flag back if bigbrain ever goes public.

Publish auth: npm Trusted Publishing (one-time setup)

CI authenticates to npm via **OIDC trusted publishing**: GitHub Actions exchanges a short-lived OIDC token for publish rights, so no long-lived `NPM_TOKEN` exists to leak. It requires two one-time setups:

1. **On npmjs.com** (a package admin, once): open the [@holokai/neuron-sdk](https://www.npmjs.com/package/@holokai/neuron-sdk) package, **Settings, Trusted Publisher, GitHub Actions**, and register:
 - Organization / repository: `holok-ai/bigbrain`
 - Workflow filename: `release.yml`
 - Environment: `npm`
2. **In this repo** (once): the `npm` GitHub Environment (Settings, Environments) with no required reviewer, because the Release run's one approval is on its cut, and a deployment branch rule of `main` only, so a copy of `release.yml` on another branch cannot publish to npm.

Until the trusted publisher is registered, the `npm publish` step fails with an auth error. That is the expected "not set up yet" signal, not a code bug.

Manual publish (fallback / first bootstrap)

The Release workflow is the norm. If you ever need to publish by hand (e.g. to bootstrap the very first version before trusted publishing is configured):

```
cd packages/neuron-sdk
pnpm run publish:npm:dry    # inspect the tarball + files list
pnpm run publish:npm       # runs the prepublishOnly build+test gate, then publishes
```

This uses your own `npm login` credentials and skips the CI approval gate — prefer the Release workflow for anything a developer will consume.